

Project: Sampling Methodologies

Purpose of the project is to compare different sampling methodologies for estimating the posterior $p(\theta|y_1, \dots, y_n)$ of a bayesian inference problem.

Data (y_1, \dots, y_n) are gene expression measurements for two genes on $n = 24$ samples where $y_i = (y_{i,1}, y_{i,2})$ represent gene expressions for sample i . Our samples are all labelled by "group", denoted (t_1, \dots, t_n) . Cell type (or mix of cell types) vary with group and we assume mean gene expression (but not variance) depends on cell type. Moreover, the expressions of different genes are independently normally distributed.

- $Y_i \sim N(\mu, \sigma^2 \mathbb{I})$ if sample i from group 1
- $Y_i \sim N(\gamma, \sigma^2 \mathbb{I})$ if sample i from group 2
- $Y_i \sim N(0.5\mu + 0.5\gamma, \sigma^2 \mathbb{I})$ if sample i from group 3
- $Y_i \sim N(\tau\mu + (1 - \tau)\gamma, \sigma^2 \mathbb{I})$ if sample i from group 4

Our model has a 6-dimensional parameter $\theta = (\sigma^2, \tau, \mu_1, \mu_2, \gamma_1, \gamma_2)$.

For all sampling methodologies, I computed 5000 samples with a burn-in of 200. The rest (step size, momentum, proposals, etc.) I discuss and treat as design choices.

1 Metropolis Hastings

To implement the Metropolis-Hastings algorithm we require the joint and transition densities to compute the hastings ratio. Below we derive the join density for our model, and treat the transition density as a design choice.

Joint density:

$$\begin{aligned}
 p(\theta|y) &\propto \prod_{i=1}^n p(y_i|\theta) p(\theta) \\
 &= \prod_{t=1}^4 \prod_{i=1}^n 1\{t_i = t\} \cdot p(y_i|\theta) \cdot p(\mu) \cdot p(\gamma) \cdot p(\tau) \cdot p(\sigma^2) \\
 &= \prod_{t_i=1} N(\mu, \sigma^2 \mathbb{I}) \cdot \prod_{t_i=2} N(\gamma, \sigma^2 \mathbb{I}) \cdot \prod_{t_i=3} N(0.5\mu + 0.5\gamma, \sigma^2 \mathbb{I}) \cdot \prod_{t_i=4} N(\tau\mu + (1 - \tau)\gamma, \sigma^2 \mathbb{I}) \cdot p(\sigma^2)
 \end{aligned}$$

Hastings Ratio: let $q(\theta_0, \theta_1)$ be our transition density and $\pi(y, \theta)$ be our joint density. Then,

$$\text{accept prob.} = \min(1, \frac{q(\theta_0, \theta_1)\pi(y, \theta_1)}{q(\theta_1, \theta_0)\pi(y, \theta_0)})$$

```

1 def metropolis_hastings(data, n_samples, step_size, initial_position):
2     curr_theta = initial_position.copy()
3     samples = [curr_theta]
4
5     while it < n_samples:
6
7         proposed = proposal_sampler(curr_theta, step_size)
8         h_ratio = hastings_ratio(proposed, curr_theta, data)
9         accept_prob = min(1, h_ratio)
10
11         if np.random.uniform(0, 1) <= accept_prob:
12             curr_theta = proposed
13             accept_count += 1
14
15         samples.append(curr_theta)
16
17     return samples

```

The Metropolis-Hastings algorithm itself is very simple, However, there are still a number of design choices we need to make, such as (a) how to propose a new sample, (b) enforcing any constraints we have (e.g. $\tau \in [0, 1]$), and (c) choosing hyperparameters (step size, burn-in).

- (a) **Enforcing constraints:** We had two constraints to consider: (a) $\sigma^2 > 0$ and (b) $0 \leq \tau \leq 1$. I enforce these constraints for our symmetric proposal method by resampling if the proposal violates either (a) or (b). This reduces the efficiency (acceptance ratio) of the algorithm substantially and introduces potential bias in our resulting marginal posterior samples (particularly for τ which has significant density near its upper bound, mean shifted from 0.75 to 0.85). By contrast, our asymmetric proposal method encodes these constraints directly. For these reasons, we use the asymmetric kernel for our results.
- (b) **Hyperparameters:** There are not many hyperparameters to tune for Metropolis-Hastings, which is one of its advantages! I experimented with different step sizes (0.01, 0.05, 0.1, 0.5), and for each computed the acceptance ratio and mean number of effective (uncorrelated) samples for 300 sample test (after burn-in). I choose a step size of 0.05 since it maximized both number of effective samples and acceptance ratio.

Step size	0.01	0.05	0.1	0.5
Acceptance ratio	0.87	0.44	0.22	0.02
Effective samples	3.5	16.6	14.6	6.96

I did not investigate varying initialization (used $[1., 0.5, 0., 0., 0., 0.]$) or burn-in (fixed at 200 to match what I saw in textbook examples).

Results: Using the above design choices, we produce the following samples from our posterior (Figure 1) and compute the first and second moments (table below). These passed basic sense checks, such as σ^2 and τ falling within their respective bounds, mean and variance matching results from other sampling algorithms, and lower variance as we increase number of samples.




Figure 1: Histogram of Metropolis-Hastings posterior samples

	σ^2	τ	μ_1	μ_2	γ_1	γ_2
Mean of posterior samples	0.14	0.84	-1.44	-0.67	-0.25	0.35
Variance of posterior samples	0.004	0.008	0.015	0.023	0.022	0.025

Final remarks. Metropolis-Hastings is simple to implement and relatively fast to run (time = 197 sec for 5000 samples), however, number of effective samples (autocorrelation) and acceptance ratio were not as good as we might expect from other techniques that take more directed paths through the posterior space (e.g. HMC).